


L03 Spatial Filtering


Introduction

- We use spatial filtering to **enhance images**, e.g., noise filtering.
- We can use the same technique to locate objects in images, called **template matching**. After completing this lesson, we would be able to process images using filters, as done in popular image processing software such as Photoshop.
- Intensity transformations affect the brightness (intensity level) of an image. The resulting value of a pixel after an intensity operation is just dependent on the original value of the same pixel.
- In contrast, the **resulting value of a pixel after a spatial filtering operation depends on the neighborhood** of the pixel in question. So, the space around the pixel in question matters, hence, the name spatial filtering.
- In linear spatial filtering we use the **2-D convolution** operation.

Spatial Filtering



A 3×3 window (kernel or filter)



Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	18	90	90	18	90	0	
3	0	0	18	90	18	18	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1									
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	

i

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30						
2									
3									
4									
5									

i

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30	40					
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30	40	40				
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		20	30	40	40	40			
2									
3									
4									
5									
i									

Spatial Filtering

	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		0	30	40	40	40	30		
2									
3									
4									
5									
i									

Spatial Filtering


	0	1	2	3	4	5	6	7	j
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	180	90	90	180	90	0	
3	0	0	180	90	180	180	0	0	
4	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	
i									

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

	0	1	2	3	4	5	6	7	j
0									
1		0	30	40	40	40	30		
2		40	60	90	90	80	50		
3		50	90	90	90	80	50		
4		20	30	50	50	40	20		
5									
i									


Practical Details: Dealing with Image Boundaries

To control the size of the output, we need to use *padding*




Practical details: Dealing with Image Boundaries

- To control the size of the output, we need to use *padding*
- What values should we pad the image with?
 - Zero pad (or clip filter)
 - Wrap around
 - Copy edge
 - Reflect across edge



Examples of Effect of Kernel Choices




Original




Averaging

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

Averaging



Sobel vertical



Sobel horizontal

-1	-2	-1
0	0	0
1	2	1

Sobel vertical

-1	0	1
-2	0	2
-1	0	1

Sobel horizontal

Averaging Using cv.filter2D

```
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt

im = cv.imread('images/sigiriya.jpg', cv.IMREAD_REDUCED_GRAYSCALE_2)
assert im is not None

kernel = np.ones((3,3),np.float32)/9
imavg = cv.filter2D(im, cv.CV_32F, kernel)

fig, axes = plt.subplots(1,2, sharex='all', sharey='all', figsize=(18,9))
axes[0].imshow(im, cmap='gray')
axes[0].set_title('Original')
axes[0].set_xticks([]), axes[0].set_yticks([])
axes[1].imshow(imavg, cmap='gray')
axes[1].set_title('Averaging')
axes[1].set_xticks([]), axes[1].set_yticks([])
plt.show()
```

Sobel Filtering Using cv.filter2D

```
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt

im = cv.imread('images/sigiriya.jpg', cv.IMREAD_REDUCED_GRAYSCALE_2)
assert im is not None

sobel_x = np.array([[ -1, -2, -1], [ 0,  0,  0], [ 1,  2,  1]])
sobel_y = np.array([[ -1,  0,  1], [-2,  0,  2], [-1,  0,  1]])

im_x = cv.filter2D(im, cv.CV_64F, sobel_x)
im_y = cv.filter2D(im, cv.CV_64F, sobel_y)

fig, ax = plt.subplots(1,2, sharex='all', sharey='all', figsize=(18,9))
ax[0].imshow(im_x, cmap='gray')
ax[0].set_title('Sobel X')
ax[0].set_xticks([]), ax[0].set_yticks([])
ax[1].imshow(im_y, cmap='gray')
ax[1].set_title('Sobel Y')
ax[1].set_xticks([]), ax[1].set_yticks([])
plt.show()
```

Convolution and Correlation

- Spatial filtering is, in fact, convolution.
- As the filters are typically symmetric—i.e., a 180° rotation results in the same kernel—correlation is equivalent to convolution.
- Correlation is also the scalar product between the kernel and the underlying image patch. Therefore, it is a measure of similarity between the kernel and the underlying image patch.
- As a result, when the kernel and the patch are “similar”, the output is high. In view of this, **spatial filtering seeks for patches in the image that are similar to the kernel.**
- Implementing filtering using loops (four nested for loops) in a non-C fashion is inefficient. Instead, use filter2D.

Convolution and Correlation

The convolution sum expression that we learned in signal and systems is

$$y[n] = x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k].$$

In 2-D, as applicable in image processing, correlation sum with a kernel $w[m, n]$ with non-zero values in $(m, n) \in ([-a, a], [-b, b])$ is

$$(w * f)[m, n] = w[m, n] * f[m, n] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t]f[m-s, n-t].$$

Correlation:

$$(w \circledast f)[m, n] = w[m, n] \circledast f[m, n] = \sum_{s=-a}^a \sum_{t=-b}^b w[s, t]f[m+s, n+t].$$

Example

Consider the image

$$f = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

and the filtering kernel

$$w = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

. Appropriately pad the image. Carry out a. correlation b. convolution.

Example

Consider the image

$$f_{\text{padded}} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad w = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

Correlation result and convolution result, respectively

$$(w \circledast f) = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 9 & 8 & 7 & 0 \\ 0 & 6 & 5 & 4 & 0 \\ 0 & 3 & 2 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad (w * f) = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 3 & 0 \\ 0 & 4 & 5 & 6 & 0 \\ 0 & 7 & 8 & 9 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Convolution: Key Properties

1. Linearity: $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$
2. Shift invariance: same behavior regardless of pixel location:
 $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$
3. Theoretical result: any linear shift-invariant operator can be represented as a convolution.

Other Properties

1. Commutative: $a * b = b * a$
2. Conceptually no difference between filter and signal
3. Associative: $a * (b * c) = (a * b) * c$
4. Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
5. This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
6. Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
7. Scalars factor out: $ka * b = a * kb = k(a * b)$
8. Identity: unit impulse $e = [\dots, 0, 0, 1, 0, 0, \dots]$, $a * e = a$

At the Edge

The filter window falls off the edge of the image.

Need to extrapolate the image.

Methods: various border types, image boundaries are denoted with '|'

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: iiiiii|abcdefgh|iiiiiii with some specified 'i'

Sharpening



(a) Original

—



(b) Smoothed

=



(c) Original – Smoothed




(d) Original

+



(e) Original – Smoothed

=



(f) Sharpened

Box Filter vs. Gaussian Filter

- What's wrong with this filtering operation?
- What's the solution?



(a) Original



(b) Kernel (much zoomed)



(c) Box Filtered

Box Filter vs. Gaussian Filter



(a) Original






(b) Box Filtered



(c) Gaussian Filtered

Gaussian Kernel

$$G_{\sigma}(x,y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$




- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case).
- The Gaussian function has infinite support, but discrete filters use finite kernels.

Choosing Gaussian Kernel Width



$$g(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Rule of thumb: set the filter half-width to about 3σ .



Separability of the Gaussian Filter

$$G(x,y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$




Note: convolution and matrix multiplication (outer product) are the same in this case

- The 2-D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y .
- In this case, the two Gaussians are the (identical) 1-D Gaussians.
- Separability means that a 2-D convolution can be reduced to two 1-D convolutions (one among rows and one among columns).
- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 $O(n^2 m^2)$.
- What is the complexity if the kernel is separable?
 $O(n^2 m)$

Gaussian Noise


- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise



$$f(x, y) = \underbrace{\hat{f}(x, y)}_{\text{Ideal Image}} + \underbrace{\eta(x, y)}_{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

$\sigma = 0.05$, no smoothing



$\sigma = 0.1$, no smoothing



$\sigma = 0.2$, no smoothing



$\sigma = 0.05$, smoothing kernel 1 px



$\sigma = 0.1$, smoothing kernel 1 px



$\sigma = 0.2$, smoothing kernel 1 px



$\sigma = 0.05$, smoothing kernel 2 px



$\sigma = 0.1$, smoothing kernel 2 px




$\sigma = 0.2$, smoothing kernel 2 px

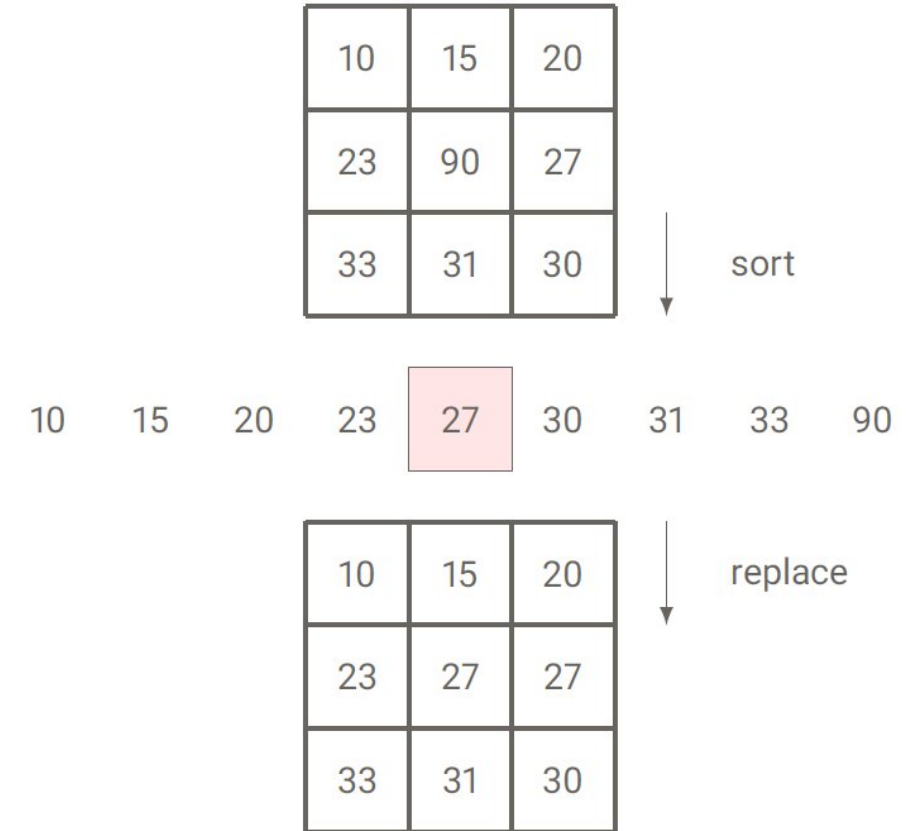


Reducing Gaussian Noise
Smoothing with larger standard deviations
suppresses noise but also blurs the image.

Reducing Salt-and-Pepper Noise: Median Filtering



Inability to reduce salt and pepper noise with Gaussian filtering



A median filter operates over a window by selecting the median intensity in the window.
Is median filtering linear?

Source: K. Grauman

Bilateral Filter

What is the Bilateral Filter?

- A non-linear, edge-preserving, and noise-reducing smoothing filter.
- Combines domain (spatial) and range (intensity) filtering.
- Useful in denoising while preserving edges, unlike Gaussian smoothing.

Mathematical Formulation

$$I_{\text{filtered}}(x) = \frac{1}{W} \sum_{p_{x_i \in \Omega}} G_{\sigma_s}(\|x_i - x\|) \cdot G_{\sigma_r}(\|I(x_i) - I(x)\|)$$

Original

Gaussian Filtered

Bilateral Filtered

